Oluwatosin Ayeni

Glasgow, Scotland, UK
07727305230 · ayenioluwatosinolawale@gmail.com
github.com/AyeniOluwatosinOlawale · linkedin.com/in/oluwatosinolawaleayeni

3 November 2026

Anthropic, London, UK

Dear Hiring Team,

I am applying for the Staff Software Engineer, Inference role at Anthropic London. I have spent two years building the class of systems this role will scale — production LLM inference platforms, GPU serving stacks, batching pipelines, KV cache optimisation, multi-provider routing, and Kubernetes-based deployment. I have done this work independently, from design through deployment to live production traffic, using Claude as my primary API provider.

Three achievements that map directly to the scope of this role:

- **Production LLM inference platform on GKE:** I built *llm-inference-platform* — a full production serving stack on Google Kubernetes Engine (NVIDIA L4 24 GB) running vLLM 0.32.0 (PagedAttention, continuous batching) and SGLang 0.5.20 (RadixAttention, scales-to-zero) behind LiteLLM Proxy 1.40.0. Gateway implements least-busy routing, A/B weighted round-robin, Redis response caching, and PostgreSQL spend tracking. KEDA autoscaling on GPU queue depth — not CPU (GPU pods at 5–15% CPU when saturated). Argo CD GitOps, Argo Rollouts canary, Traefik TLS, Terraform IaC. LIVE: 723 requests, 1.25M+ tokens, $152.82 over 5.3 days.
- **H100 inference benchmarking — published results:** I benchmarked vLLM v0.30 vs SGLang v0.5.20 on 2× NVIDIA H100 NVL 94 GB (NVLink 4.0, CUDA 13.4), Qwen3-8B BF16, TP=2, across concurrency 1–128. Peak: 5,965 tok/s. RadixAttention: 1,443× TTFT speedup (39 ms vs 55,565 ms) at 131K-token context. 100% SLO compliance. Results live at inference-web-theta.vercel.app/benchmarks/llm.
- **Inference Gateway — routing, caching, batching at scale:** I built *inference-optimizer* — a production AI Inference Gateway implementing 13 optimisation strategies across Claude, GPT-4o, and Gemini: parallel execution (3.87× faster), prompt caching (61.4% cost reduction), model routing (64.3% cheaper), and dynamic concurrency management (14× throughput).

The London team working on Claude's inference infrastructure is exactly where I want to apply this experience. I use Claude in production daily and I take the safety-reliability connection seriously. I am planning to relocate to London and available to interview at your convenience.

Yours sincerely,

Oluwatosin Ayeni

github.com/AyeniOluwatosinOlawale
linkedin.com/in/oluwatosinolawaleayeni