Oluwatosin Ayeni

Glasgow, Scotland, UK
07727305230 · ayenioluwatosinolawale@gmail.com
github.com/AyeniOluwatosinOlawale · linkedin.com/in/oluwatosinolawaleayeni

28 October 2026

Leonardo UK, Edinburgh

Dear Hiring Team,

I am applying for the AI Platform Engineer role at Leonardo UK. The platform layer you are describing — model serving, inference runtimes, AI gateways, and Kubernetes-based GPU infrastructure — is precisely what I have been building independently, and the Edinburgh hybrid arrangement is ideal from my base in Glasgow.

Three achievements that map directly to what this role requires:

- **Production Kubernetes LLM serving platform:** I built *vllm-k8s-stack* — production LLM inference on Kubernetes using a vLLM serving engine, KV-cache-aware request routing, HPA autoscaling on GPU node pools, and a Prometheus/Grafana observability stack. This is the same platform layer above managed GPU infrastructure that Leonardo is looking to own and operate — designed for multi-tenant concurrent workloads with resource controls and availability monitoring. github.com/AyeniOluwatosinOlawale/vllm-k8s-stack
- **Production AI Gateway — authentication, routing, rate limiting:** I built *inference-optimizer*, a production AI Inference Gateway spanning Claude, GPT-4o, Gemini, and self-hosted vLLM. It handles authentication, model routing (64.3% cost reduction), rate limiting, parallel execution (3.87× faster), and prompt caching (61.4% cost reduction) — with real-time usage monitoring and audit logging. This is the practical equivalent of LiteLLM or Bifrost deployed as an enterprise AI gateway, which Leonardo explicitly lists as a key area of experience. Live at inference-web-theta.vercel.app.
- **H100 inference platform operations:** I benchmarked and operated vLLM and SGLang on 2× NVIDIA H100 NVL 94 GB hardware — measuring throughput (5,965 tok/s peak), TTFT, TPOT, and SLO compliance across concurrency 1–128. 100% SLO compliance. I also profiled GPU utilisation with NVIDIA Nsight Systems to identify and resolve performance bottlenecks — the same GPU performance and utilisation management this role requires.

Leonardo's focus on making AI capabilities easy to consume for engineering teams while maintaining security, performance, and availability controls resonates with how I have approached building internal AI platform infrastructure. Operating in a secure enterprise environment with restricted internet access is a constraint I have designed around directly — API credential management, offline-capable service design, and air-gap-friendly deployment patterns are part of my operational practice.

I am happy to discuss the architecture of any of these systems in detail, walk through the Kubernetes serving configuration, or demonstrate the gateway observability setup. I look forward to hearing from you.

Yours sincerely,

Oluwatosin Ayeni

github.com/AyeniOluwatosinOlawale
linkedin.com/in/oluwatosinolawaleayeni