Oluwatosin Ayeni

Glasgow, Scotland, UK
07727305230 · ayenioluwatosinolawale@gmail.com
github.com/AyeniOluwatosinOlawale · linkedin.com/in/oluwatosinolawaleayeni

4 November 2026

Nebius

Dear Hiring Team,

I am writing to apply for the ML Engineer, LLM Training & Inference Optimization role at Nebius. Nebius's applied AI platform and Token Factory's focus on efficient open model deployment is exactly where I want to contribute.

Three specific achievements directly relevant to this role:

- **H100 NVL inference:** I benchmarked vLLM v0.30 vs SGLang v0.5.20 on 2× NVIDIA H100 NVL 94 GB (NVLink 4.0, CUDA 13.4), Qwen3-8B BF16, TP=2. RadixAttention prefix caching: 1,443× TTFT speedup (39 ms vs 55,565 ms) at 131K-token context; 5,965 tok/s peak; 100% SLO compliance at concurrency 1–128. Live: inference-web-theta.vercel.app/benchmarks/llm
- **Production cost and throughput:** Via vLLM/SGLang continuous batching and KV cache optimisation on RTX 4090, I achieved 28× LLM cost reduction and 3.8× throughput improvement. CUDA KV cache kernel profiling with Nsight identified memory bandwidth as bottleneck; 18× latency reduction via continuous batching redesign.
- **Full LLM lifecycle:** Fine-tuning pipelines (Mistral 7B, Phi-4, Qwen3 via LoRA/PEFT), trained LM from scratch (Qwen3-TinyStories), AWQ/GPTQ quantisation, model architecture labs (MHA/GQA/MLA, Dense vs MoE, VRAM modelling).

I look forward to hearing from you.

Yours sincerely,

Oluwatosin Ayeni

github.com/AyeniOluwatosinOlawale
linkedin.com/in/oluwatosinolawaleayeni