Oluwatosin Ayeni

AI Platform Engineer — Leonardo UK

Glasgow, Scotland, UK · 07727305230 · ayenioluwatosinolawale@gmail.com
github.com/AyeniOluwatosinOlawale · linkedin.com/in/oluwatosinolawaleayeni

PROFILE

AI Platform Engineer with production experience deploying and operating LLM inference services on Kubernetes, building AI gateway infrastructure, and managing GPU workloads at scale. Benchmarked vLLM and SGLang on H100 NVL achieving 5,965 tok/s peak and 100% SLO compliance; built KV-cache-aware routing, HPA autoscaling, and Prometheus/Grafana observability stacks. Experienced designing multi-provider AI gateways with authentication, rate limiting, and model routing. Based in Glasgow — Edinburgh hybrid is ideal.

CORE COMPETENCIES

Inference & Serving	vLLM, SGLang, TensorRT-LLM, LiteLLM-style AI gateways, model routing, rate limiting, multi-provider API management
Containers & Orchestration	Kubernetes (HPA, GPU node pools, resource quotas), Docker, Helm, container networking, pod security policies
GPU & Performance	NVIDIA Nsight Systems/Compute, GPU utilisation monitoring, CUDA kernel profiling, tensor parallelism, KV cache optimisation, BFloat16/FP8
Observability & SRE	Prometheus, Grafana, distributed tracing, SLO/SLA compliance monitoring, alerting, incident response
Platform Engineering	Python, Bash, FastAPI, PostgreSQL, Redis, GitHub Actions CI/CD, infrastructure as code, secure service design, authentication/access control

PROFESSIONAL EXPERIENCE

Jan 2025–present — **Founder & Lead Engineer** — *Nexora Academic*, Glasgow, UK

- H100 NVL benchmark: vLLM v0.30 vs SGLang v0.5.20 — Qwen3-8B BF16, TP=2. Peak: 5,965 tok/s; RadixAttention: 1,443× TTFT speedup (39 ms vs 55,565 ms) at 131K-token context; 100% SLO compliance up to concurrency 128. Live: inference-web-theta.vercel.app
- Built *inference-optimizer*: production AI Inference Gateway — 3.87× faster (parallel execution), 61.4% cost reduction (prompt caching), 14× throughput (concurrency), 64.3% cheaper (model routing). 13 optimisations, 4 providers.
- Built *vllm-k8s-stack*: Kubernetes LLM serving — KV-cache-aware routing, HPA autoscaling on GPU pools, Prometheus/Grafana observability.
- Built *inference-harness*: unified benchmark framework spanning closed APIs + vLLM/SGLang/TensorRT-LLM with GPU profiling and live dashboard.
- Designed 13-phase AI Inference Engineering curriculum covering GPU architecture, CUDA optimisation, serving frameworks, and distributed inference.

Jan 2023–present — **Founder & AI Consultant** — *Vertinab / AnaburyAI Technology*, Glasgow, UK

- 28× LLM cost reduction deploying Qwen2.5-7B on RTX 4090 via vLLM/SGLang with continuous batching and KV cache optimisation.
- CUDA KV cache kernel profiling with NVIDIA Nsight: identified memory bandwidth bottleneck; 3.8× throughput improvement; 18× latency reduction via continuous batching redesign.
- Fine-tuned Mistral 7B with LoRA/PEFT on radiology corpus: +22% diagnostic accuracy, −35% hallucination rate; deployed as production clinical inference API.

2024–present — **Senior Software Engineer (Contract)** — *EduTechConnect / Stellar Labs*, Remote

- Real-time NLP pipeline (WhatsApp → Zoho CRM → Supabase) with sub-500 ms end-to-end latency under concurrent load.

Mar 2026–present **AI Automation Engineer** — *EduTech Global*, Remote

- LLM-backed student chatbots with structured output extraction and multi-step reasoning for Ahmadu Bello and Babcock universities.

2018–2022 **Data Analyst** — *Bytelab Technologies Ltd*, Nigeria

KEY PROJECTS

vllm-k8s-stack Production Kubernetes LLM serving: vLLM, KV-cache-aware routing, HPA GPU autoscaling, Prometheus/Grafana observability. GitHub

inference-optimizer Production AI Gateway: 4 providers, authentication, model routing, rate limiting, caching, monitoring. 3.87× faster, 14× throughput, 61.4% cost reduction. GitHub

llm-inference-benchmark (H100) vLLM vs SGLang, H100 NVL — throughput, TTFT, TPOT, SLO compliance at concurrency 1–128. 100% SLO. Live dashboard

inference-harness Unified benchmark harness: closed + open APIs, GPU profiling, automated reporting. GitHub

EDUCATION

2022–2023 **MSc Artificial Intelligence and Applications, Distinction** — *University of Strathclyde*, Glasgow, UK

2010–2016 **MTech / BSc Physics** — *Federal University of Technology Akure*, Nigeria

OPEN SOURCE

GitHub 43 public repos — github.com/AyeniOluwatosinOlawale